

E-SMI Library: EPYC™ Systems Management Interface Library

Generated by Doxygen 1.8.14

Contents

1	EPYC™ System Management Interface (E-SMI) In-band Library	1
2	Module Index	9
2.1	Modules	9
3	Data Structure Index	11
3.1	Data Structures	11
4	File Index	13
4.1	File List	13
5	Module Documentation	15
5.1	Initialization and Shutdown	15
5.1.1	Detailed Description	15
5.1.2	Function Documentation	15
5.1.2.1	esmi_init()	15
5.2	Energy Monitor (RAPL MSR)	16
5.2.1	Detailed Description	16
5.2.2	Function Documentation	16
5.2.2.1	esmi_core_energy_get()	16
5.2.2.2	esmi_socket_energy_get()	17
5.2.2.3	esmi_all_energies_get()	17
5.3	HSMP System Statistics	18
5.3.1	Detailed Description	18
5.3.2	Function Documentation	18

5.3.2.1	esmi_smu_fw_version_get()	18
5.3.2.2	esmi_prochot_status_get()	19
5.3.2.3	esmi_fclk_mclk_get()	19
5.3.2.4	esmi_cclk_limit_get()	20
5.3.2.5	esmi_hsmp_proto_ver_get()	20
5.3.2.6	esmi_socket_current_active_freq_limit_get()	20
5.3.2.7	esmi_socket_freq_range_get()	21
5.3.2.8	esmi_current_freq_limit_core_get()	21
5.4	Power Monitor	23
5.4.1	Detailed Description	23
5.4.2	Function Documentation	23
5.4.2.1	esmi_socket_power_get()	23
5.4.2.2	esmi_socket_power_cap_get()	23
5.4.2.3	esmi_socket_power_cap_max_get()	24
5.4.2.4	esmi_pwr_svi_telemetry_all_rails_get()	24
5.5	Power Control	27
5.5.1	Detailed Description	27
5.5.2	Function Documentation	27
5.5.2.1	esmi_socket_power_cap_set()	27
5.5.2.2	esmi_pwr_efficiency_mode_set()	28
5.6	Performance (Boost limit) Monitor	29
5.6.1	Detailed Description	29
5.6.2	Function Documentation	29
5.6.2.1	esmi_core_boostlimit_get()	29
5.6.2.2	esmi_socket_c0_residency_get()	29
5.7	Performance (Boost limit) Control	31
5.7.1	Detailed Description	31
5.7.2	Function Documentation	31
5.7.2.1	esmi_core_boostlimit_set()	31
5.7.2.2	esmi_socket_boostlimit_set()	32

5.8	ddr_bandwidth Monitor	33
5.8.1	Detailed Description	33
5.8.2	Function Documentation	33
5.8.2.1	esmi_ddr_bw_get()	33
5.9	Temperature Query	34
5.9.1	Detailed Description	34
5.9.2	Function Documentation	34
5.9.2.1	esmi_socket_temperature_get()	34
5.10	Dimm statistics	35
5.10.1	Detailed Description	35
5.10.2	Function Documentation	35
5.10.2.1	esmi_dimm_temp_range_and_refresh_rate_get()	35
5.10.2.2	esmi_dimm_power_consumption_get()	36
5.10.2.3	esmi_dimm_thermal_sensor_get()	36
5.11	xGMI bandwidth control	37
5.11.1	Detailed Description	37
5.11.2	Function Documentation	37
5.11.2.1	esmi_xgmi_width_set()	37
5.12	GMI3 width control	38
5.12.1	Detailed Description	38
5.12.2	Function Documentation	38
5.12.2.1	esmi_gmi3_link_width_range_set()	38
5.13	APB and LCLK level control	39
5.13.1	Detailed Description	39
5.13.2	Function Documentation	39
5.13.2.1	esmi_apb_enable()	39
5.13.2.2	esmi_apb_disable()	40
5.13.2.3	esmi_socket_lclk_dpm_level_set()	40
5.13.2.4	esmi_socket_lclk_dpm_level_get()	41
5.13.2.5	esmi_pcie_link_rate_set()	41

5.13.2.6	<code>esmi_df_pstate_range_set()</code>	42
5.14	Bandwidth Monitor	43
5.14.1	Detailed Description	43
5.14.2	Function Documentation	43
5.14.2.1	<code>esmi_current_io_bandwidth_get()</code>	43
5.14.2.2	<code>esmi_current_xgmi_bw_get()</code>	43
5.15	Metrics Table	45
5.15.1	Detailed Description	45
5.15.2	Function Documentation	45
5.15.2.1	<code>esmi_metrics_table_version_get()</code>	45
5.15.2.2	<code>esmi_metrics_table_get()</code>	45
5.15.2.3	<code>esmi_dram_address_metrics_table_get()</code>	47
5.16	Test HSMP mailbox	48
5.16.1	Detailed Description	48
5.16.2	Function Documentation	48
5.16.2.1	<code>esmi_test_hsmp_mailbox()</code>	48
5.17	Auxiliary functions	49
5.17.1	Detailed Description	49
5.17.2	Function Documentation	49
5.17.2.1	<code>esmi_cpu_family_get()</code>	49
5.17.2.2	<code>esmi_cpu_model_get()</code>	50
5.17.2.3	<code>esmi_threads_per_core_get()</code>	50
5.17.2.4	<code>esmi_number_of_cpus_get()</code>	50
5.17.2.5	<code>esmi_number_of_sockets_get()</code>	51
5.17.2.6	<code>esmi_first_online_core_on_socket()</code>	51
5.17.2.7	<code>esmi_get_err_msg()</code>	52

6	Data Structure Documentation	53
6.1	ddr_bw_metrics Struct Reference	53
6.1.1	Detailed Description	53
6.2	dimm_power Struct Reference	53
6.2.1	Detailed Description	54
6.3	dimm_thermal Struct Reference	54
6.3.1	Detailed Description	54
6.4	dpm_level Struct Reference	54
6.4.1	Detailed Description	55
6.5	link_id_bw_type Struct Reference	55
6.5.1	Detailed Description	55
6.6	smu_fw_version Struct Reference	55
6.6.1	Detailed Description	56
6.7	temp_range_refresh_rate Struct Reference	56
6.7.1	Detailed Description	56
7	File Documentation	57
7.1	e_smi.h File Reference	57
7.1.1	Detailed Description	60
7.1.2	Enumeration Type Documentation	60
7.1.2.1	io_bw_encoding	60
7.1.2.2	esmi_status_t	60
	Index	63

Chapter 1

EPYC™ System Management Interface (E-SMI) In-band Library

The EPYC™ System Management Interface In-band Library, or E-SMI library, is part of the EPYC™ System Management Inband software stack. It is a C library for Linux that provides a user space interface to monitor and control the CPU's power, energy, performance and other system management features.

Important note about Versioning and Backward Compatibility

The E-SMI library is currently under development, and therefore subject to change at the API level. The intention is to keep the API as stable as possible while in development, but in some cases we may need to break backwards compatibility in order to achieve future stability and usability. Following Semantic Versioning rules, while the E-SMI library is in a high state of change, the major version will remain 0, and achieving backward compatibility may not be possible.

Once new development has leveled off, the major version will become greater than 0, and backward compatibility will be enforced between major versions.

Building E-SMI

Downloading the source

The source code for E-SMI library is available at [Github](#).

Directory structure of the source

Once the E-SMI library source has been cloned to a local Linux machine, the directory structure of source is as below:

- `$ docs/` Contains Doxygen configuration files and Library descriptions
- `$ tools/` Contains e-smi tool, based on the E-SMI library
- `$ include/` Contains the header files used by the E-SMI library
- `$ src/` Contains library E-SMI source
- `$ cmake_modules/` Contains helper utilities for determining package and library version
- `$ DEBIAN/` Contains debian pre and post installation scripts
- `$ RPM/` Contains rpm pre and post installation scripts

Building the library and tool

Building the library is achieved by following the typical CMake build sequence, as below

- `$ cd <location of root of E-smi library>`
- `$ mkdir -p build`
- `$ cd build`
- `$ cmake ../`

Building the library for static linking

Building the library as a static(.a) along with shared libraries(.so) is achieved by following sequence. The static library is part of RPM and DEB package when compiled with cmake as below and built with 'make package'.

- `$ cmake -DENABLE_STATIC_LIB=1 ../`
- `$ make`
The built library `libe_smi64_static.a`, `libe_smi64.so.X.Y` and `esmi_tool` will appear in the build directory
- `$ sudo make install`
Library file, header and tool are installed at `/opt/e-smi`

Note: Library is dependent on `amd_hsmpl.h` header and without this, compilation will break. Please follow the instruction in "Kernel dependencies" section

Building the Documentation

The documentation PDF file can be built with the following steps (continued from the steps above) `$ make doc`
Upon a successful build, the `ESMI_Manual.pdf` and `ESMI_IB_Release_Notes.pdf` will be copied to the top directory of the source.

Building the package

The RPM and DEB packages can be created with the following steps (continued from the steps above): `$ make package`

Kernel dependencies

The E-SMI Library depends on the following device drivers from Linux to manage the system management features.

Monitoring energy counters

- AMD family 19, model 00-0fh and 30-3fh
 - These processors support energy monitoring through 32 bit RAPL MSR registers. `*amd_energy` driver, an out of tree kernel module, hosted at `amd_energy` can report per core and per socket counters via the HWMON sysfs entries.
 - This driver provides accumulation of energy for avoiding wrap around problem.
- AMD family 19, model 10-1fh and a0-afh
 - These processors support energy monitoring through 64 bit RAPL MSR registers.
 - Because of 64 bit registers, there is no accumulation of energy needed.
 - For these processors "msr-safe" driver is used from `msr-safe`.
 - Msr-safe driver needs allowlist file to be written to `/dev/cpu/msr_allowlist` for allowing the read of those specific msr registers. Please follow below steps or use `–writeallowlist` tool option.
 - * create "amd_allowlist" file with below contents and run the command "sudo su" and "cat amd_↵
allowlist > /dev/cpu/msr_allowlist"

```
# MSR # Write Mask # Comment
0xC0010299 0x0000000000000000 # "ENERGY_PWR_UNIT_MSR"
0xC001029A 0x0000000000000000 # "ENERGY_CORE_MSR"
0xC001029B 0x0000000000000000 # "ENERGY_PKG_MSR"
```

Monitoring and managing power metrics, boostlimits and other system management features

The power metrics, boostlimits and other features are managed by the SMU(System Management Unit of the processor) firmware and exposed via PCI config space and accessed through "Host System Management Port(H↵SMP)" at host/cpu side. AMD provides Linux kernel module(`amd_hsm`) exposing this information to the user-space via `ioctl` interface.

- `amd_hsm` driver is accepted in upstream kernel and is available at linux tree at `drivers/platform/x86/amd/hsm`.↵
c from version 5.17.rc1 onwards
- If you are using a kernel version less than that, then copy the header file from `arch/x86/include/uapi/asm/amd↵_hsm.h` in linux source tree to below locations based on the type of the system.
 - On RHEL systems, path is: `/usr/include/asm/`
 - On Ubuntu systems path is: `/usr/include/x86_64-linux-gnu/asm/`
- To get HSMP working. PCIe interface needs to be enabled in the BIOS. On the reference BIOS please follow the sequence below for enabling HSMP.

Advanced > AMD CBS > NBIO Common Options > SMU Common Options > HSMP Support

BIOS Default: "Auto" (Disabled) to BIOS Default: "Enabled"

If the above HSMP support option is disabled, the related E-SMI APIs will return -ETIMEDOUT.

Supported hardware

AMD Zen3 based CPU Family 19h Models 0h-Fh and 30h-3Fh.

AMD Zen4 based CPU Family 19h Models 10h-1Fh and A0-AFh.

Additional required software for building

In order to build the E-SMI library, the following components are required. Note that the software versions listed are what is being used in development. Earlier versions are not guaranteed to work:

- CMake (v3.5.0)

In order to build the latest documentation, the following are required:

- DOxygen (1.8.13)
- latex (pdfTeX 3.14159265-2.6-1.40.18)

Library Usage Basics

Many of the functions in the library take a "core/socket index". The core/socket index is a number greater than or equal to 0, and less than the number of cores/sockets on the system. Number of cores/sockets in a system can be obtained from esmi library APIs.

Hello E-SMI

The only required E-SMI call for any program that wants to use E-SMI is the `esmi_init()` call. This call initializes some internal data structures that will be used by subsequent E-SMI calls.

When E-SMI is no longer being used, `esmi_exit()` should be called. This provides a way to do any releasing of resources that E-SMI may have held. In many cases, this may have no effect, but may be necessary in future versions of the library.

Below is a simple "Hello World" type program that display the Average Power of Sockets.

```
#include <stdio.h>
#include <stdint.h>
#include <e_smi/e_smi.h>
#include <e_smi/e_smi_monitor.h>

int main()
{
    esmi_status_t ret;
    unsigned int i;
    uint32_t power;
    uint32_t total_sockets = 0;

    ret = esmi_init();
    if (ret != ESMI_SUCCESS) {
        printf("ESMI Not initialized, drivers not found.\n"
            "Err[%d]: %s\n", ret, esmi_get_err_msg(ret));
        return ret;
    }

    total_sockets = esmi_get_number_of_sockets();
    for (i = 0; i < total_sockets; i++) {
        power = 0;
        ret = esmi_socket_power_get(i, &power);
        if (ret != ESMI_SUCCESS) {
            printf("Failed to get socket[%d] avg_power, "
                "Err[%d]: %s\n", i, ret, esmi_get_err_msg(ret));
        }
        printf("socket_%d_avgpower = %.3f Watts\n",
            i, (double)power/1000);
    }
    esmi_exit();

    return ret;
}
```

Tool Usage

E-SMI tool is a C program based on the E-SMI In-band Library, the executable "e_smi_tool" will be generated in the build/ folder. This tool provides options to Monitor and Control System Management functionality.

Below is a sample usage to dump core and socket metrics

```
e_smi_library/b$ sudo ./e_smi_tool
```

```
===== E-SMI =====
```

```
-----
| CPU Family      | 0x19 (25 ) |
| CPU Model      | 0x10 (16 ) |
| NR_CPUS        | 384        |
| NR_SOCKETS     | 2          |
| THREADS PER CORE | 2 (SMT ON) |
-----
```

```
-----
| Sensor Name          | Socket 0      | Socket 1      |
-----
| Energy (K Joules)    | 14437.971     | 14087.151     |
| Power (Watts)        | 174.290       | 169.630       |
| PowerLimit (Watts)   | 400.000       | 320.000       |
| PowerLimitMax (Watts)| 400.000       | 320.000       |
| C0 Residency (%)     | 0             | 0             |
| DDR Bandwidth        |               |               |
|   DDR Max BW (GB/s)  | 58            | 58            |
|   DDR Utilized BW (GB/s) | 0            | 0            |
|   DDR Utilized Percent(%) | 0            | 0            |
| Current Active Freq limit |             |             |
|   Freq limit (MHz)   | 3500          | 3500          |
|   Freq limit source  | Refer below[*0] | Refer below[*1] |
| Socket frequency range |             |             |
|   Fmax (MHz)         | 3500          | 3500          |
|   Fmin (MHz)         | 400           | 400           |
-----
```

```
-----
| CPU energies in Joules:
| cpu [ 0 ] :      645.992   181.415   171.678   165.577   161.001   158.397   161.333   151.716
| cpu [ 8 ] :      88.197    79.306    73.860    73.015    72.960    69.293    67.871    78.895
| cpu [ 16 ] :     70.376    71.231    61.756    63.061    80.656    73.360    69.566    69.969
| cpu [ 24 ] :     67.054    65.621    64.468    66.346    64.344    64.310    71.548    65.579
| cpu [ 32 ] :     65.731    62.931    65.526    69.765    69.050    65.782    70.630    65.282
| cpu [ 40 ] :     69.608    67.261    63.765    69.477    68.677    63.145    62.451    159.949
| cpu [ 48 ] :     70.810    73.084    64.584    62.966    66.581    65.620    62.381    65.602
| cpu [ 56 ] :     72.804    70.842    69.651    64.990    63.924    66.468    63.401    296.924
| cpu [ 64 ] :     64.693    62.723    65.057    62.515    60.091    60.422    62.217    66.552
| cpu [ 72 ] :     81.746    70.622    68.848    301.949    78.974    68.130    68.141    65.693
| cpu [ 80 ] :     77.475    72.441    81.296    71.441    71.988    75.237    73.986    69.467
| cpu [ 88 ] :     73.385    69.277    61.759    61.060    62.834    60.681    62.835    62.703
| cpu [ 96 ] :    142.718   139.519   134.449   134.097   135.045   140.307   140.553   137.153
| cpu [104] :     66.016    66.736    62.224    67.137    64.881    70.592    64.701    64.056
| cpu [112] :     70.791    69.107    70.638    69.998    68.199    65.263    70.638    72.557
| cpu [120] :     94.391    94.151    71.881    66.493    64.653    66.141    66.132    69.593
| cpu [128] :     65.800    64.742    63.130    61.771    65.416    66.205    64.663    71.349
| cpu [136] :     72.183    66.754    67.090    63.343    69.450    67.979    68.285    70.478
| cpu [144] :     68.281    63.809    62.717    63.348    71.164    72.289    65.516    65.513
| cpu [152] :     74.588    69.074    66.711    66.011    67.896    65.933    67.031    65.474
-----
```

```

| cpu [160] :      66.668      62.996      65.945      63.734      64.060      68.597      76.405      91.436
| cpu [168] :      77.658      70.085      67.025      68.951      64.678      64.821      65.031      71.694
| cpu [176] :      72.782      89.196      74.777      73.703      66.247      65.419      64.748      63.978
| cpu [184] :      63.887      66.080      64.042      65.151      69.661      74.616      63.834      69.824

```

```

-----
| CPU boostlimit in MHz:

```

```

| cpu [ 0] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [16] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [32] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [48] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [64] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [80] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [96] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [112] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [128] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [144] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [160] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [176] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500

```

```

-----
| CPU core clock current frequency limit in MHz:

```

```

| cpu [ 0] : 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500 3500
| cpu [16] : NA   NA   NA   NA   NA   NA   NA   NA   3500 3500 3500 3500 3500 3500 3500
| cpu [32] : 3500 3500 3500 3500 3500 3500 3500 3500 NA   NA   NA   NA   NA   NA   NA
| cpu [48] : 3500 3500 3500 3500 3500 3500 3500 3500 NA   NA   NA   NA   NA   NA   NA
| cpu [64] : NA   NA   NA   NA   NA   NA   NA   NA   3500 3500 3500 3500 3500 3500 3500
| cpu [80] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [96] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [112] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [128] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [144] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [160] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA
| cpu [176] : NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA   NA

```

```

-----
*0 Frequency limit source names:
OPN Max

```

```

*1 Frequency limit source names:
OPN Max

```

Try './e_smi_tool --help' for more information.

For detailed and up to date usage information, we recommend consulting the help:

For convenience purposes, following is the output from the -h flag:

```
e_smi_library/b$ ./e_smi_tool -h
```

```
===== E-SMI =====
```

```
Usage: ./e_smi_tool [Option]... <INPUT>...
```

```
Output Option<s>:
```

```
-h, --help           Show this help message
-A, --showall        Show all esmi parameter values
-V, --version        Show e-smi library version
--testmailbox [SOCKET] [VALUE] Test HSMP mailbox interface
--writemsrallowlist  Write msr-safe allowlist file
```

```
Get Option<s>:
```

```
--showcoreenergy [CORE]           Show energy for a given CPU (Joules)
--showsockenergy                Show energy for all sockets (KJoules)
--showsockpower                 Show power metrics for all sockets (Watts)
--showcorebl [CORE]             Show Boostlimit for a given CPU (MHz)
--showsockcOres [SOCKET]        Show cO_residency for a given socket (%)
--showsmufwver                  Show SMU FW Version
--showhsmpprotover              Show HSMP Protocol Version
--showprochotstatus             Show HSMP PROCHOT status for all sockets
--showclocks                    Show Clock Metrics (MHz) for all sockets
--showcclkfreqlimit [CORE]      Show current clock frequency limit(MHz) for a
    given core
--showsvipower                  Show svi based power telemetry of all rails
    for all sockets
--showxgmibw [LINK<P2,P3,G0-G7>] [BW<AGG_BW,RD_BW,WR_BW>] Show xGMI bandwidth for a given socket,
    linkname and bwtype
--showiobw [SOCKET] [LINK<P2,P3,G0-G7>] Show IO aggregate bandwidth for a given
    socket and linkname
--showlclkdpmlvl [SOCKET] [NBIOID<0-3>] Show lclk dpm level for a given nbio in a
    given socket
--showsockcclkfreqlimit [SOCKET] Show current clock frequency limit(MHz) for a
    given socket
--showmetrictablever            Show Metrics Table Version
--showmetrictable [SOCKET]     Show Metrics Table

--setpowerlimit [SOCKET] [POWER] Set power limit for a given socket (mWatts)
--setcorebl [CORE] [BOOSTLIMIT] Set boost limit for a given core (MHz)
--setsockbl [SOCKET] [BOOSTLIMIT] Set Boost limit for a given Socket (MHz)
--setxgmewidth [MIN<0-2>] [MAX<0-2>] Set xgmi link width in a multi socket system
    (MAX >= MIN)
--setlclkdpmlvl [SOCKET] [NBIOID<0-3>] [MIN<0-3>] [MAX<0-3>] Set lclk dpm level for a given nbio in a
    given socket (MAX >= MIN)
```

```
===== End of E-SMI =====
```

Following are the value ranges and other information needed for passing it to tool

1. `--showxgmibw [SOCKET] [LINKNAME] [BWTYPE]`
 LINKNAME :
 Rolling Stones:P0/P1/P2/P3/G0/G1/G2/G3
 Mi300A:G0/G1/G2/G3/G4/G5/G6/G7
 Breithorn:P1/P3/G0/G1/G2/G3
 BWTYPE : AGG_BW/RD_BW/WR_BW
2. `--setxgmewidth [MIN] [MAX]`
 MIN : MAX : 0 - 2 with MIN <= MAX
3. `--showlclkdpmlvl [SOCKET] [NBIOID]`
 NBIOID : 0 - 3
4. `--apbdisable [SOCKET] [PSTATE]`
 PSTATE : 0 - 2
5. `--setlclkdpmlvl [SOCKET] [NBIOID] [MIN] [MAX]`
 NBIOID : 0 - 3
 MIN : MAX : 0 - 3 with MIN <= MAX
6. `--setpcielinekratecontrol [SOCKET] [CTL]`
 CTL : 0 - 2
7. `--setpowerefficiencymode [SOCKET] [MODE]`

```

MODE : 0 - 3

8. --setdfpstatetrange [SOCKET] [MAX] [MIN]

    MIN : MAX : 0 - 2 with MAX <= MIN

9. --setgmi3linkwidth [SOCKET] [MIN] [MAX]

    MIN : MAX : 0 - 2 with MIN <= MAX

10. --testmailbox [SOCKET] [VALUE]
    VALUE : Any 32 bit value

```

Below is a sample usage to get different system metrics information

```

1. e_smi_library/b$ sudo ./e_smi_tool --showcoreenergy 0

===== E-SMI =====

-----
| core[000] energy |           646.549 Joules |
-----

===== End of E-SMI =====

2. e_smi_library/b$ sudo ./e_smi_tool --showcoreenergy 12 --showsockpower --setpowerlimit 1 220000
   --showsockpower

===== E-SMI =====

-----
| core[012] energy |           73.467 Joules |
-----

-----
| Sensor Name           | Socket 0           | Socket 1           |
-----
| Power (Watts)         | 174.051            | 169.451            |
| PowerLimit (Watts)    | 400.000            | 220.000            |
| PowerLimitMax (Watts) | 400.000            | 320.000            |
-----

Socket[1] power_limit set to 220.000 Watts successfully

-----
| Sensor Name           | Socket 0           | Socket 1           |
-----
| Power (Watts)         | 174.085            | 169.431            |
| PowerLimit (Watts)    | 400.000            | 220.000            |
| PowerLimitMax (Watts) | 400.000            | 320.000            |
-----

3. e_smi_library/b$ ./e_smi_tool --showxgmibandwidth G2 AGG_BW

===== E-SMI =====

-----
| Current Aggregate bandwidth of xGMI link G2 |    40 Mbps |
-----

===== End of E-SMI =====

4. e_smi_library/b$ sudo ./e_smi_tool --setdfpstatetrange 0 1 2
[sudo] password for user:

===== E-SMI =====

Data Fabric PState range(max:1 min:2) set successfully

===== End of E-SMI =====

```


Chapter 2

Module Index

2.1 Modules

Here is a list of all modules:

Initialization and Shutdown	15
Energy Monitor (RAPL MSR)	16
HSMP System Statistics	18
Power Monitor	23
Power Control	27
Performance (Boost limit) Monitor	29
Performance (Boost limit) Control	31
ddr_bandwidth Monitor	33
Temperature Query	34
Dimm statistics	35
xGMI bandwidth control	37
GMI3 width control	38
APB and LCLK level control	39
Bandwidth Monitor	43
Metrics Table	45
Test HSMP mailbox	48
Auxiliary functions	49

Chapter 3

Data Structure Index

3.1 Data Structures

Here are the data structures with brief descriptions:

ddr_bw_metrics	DDR bandwidth metrics	53
dimm_power	DIMM Power(mW), power update rate(ms) and dimm address	53
dimm_thermal	DIMM temperature(°C) and update rate(ms) and dimm address	54
dpm_level	Max and min LCLK DPM level on a given NBIO ID. Valid max and min DPM level values are 0 - 1	54
link_id_bw_type	LINK name and Bandwidth type Information.It contains link names i.e valid link names are "↔ P0", "P1", "P2", "P3", "P4", "G0", "G1", "G2", "G3", "G4" "G5", "G6", "G7" Valid bandwidth types 1(Aggregate_BW), 2 (Read BW), 4 (Write BW)	55
smu_fw_version	Deconstruct raw uint32_t into SMU firmware major and minor version numbers	55
temp_range_refresh_rate	Temperature range and refresh rate metrics of a DIMM	56

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

e_smi.h	57
-----------------------------------	----

Chapter 5

Module Documentation

5.1 Initialization and Shutdown

Functions

- `esmi_status_t esmi_init` (void)
Initialize the library, validates the dependencies.
- `void esmi_exit` (void)
Clean up any allocation done during init.

5.1.1 Detailed Description

This function validates the dependencies that exist and initializes the library.

5.1.2 Function Documentation

5.1.2.1 `esmi_init()`

```
esmi_status_t esmi_init (  
    void )
```

Initialize the library, validates the dependencies.

Search the available dependency entries and initialize the library accordingly.

Return values

<code>ESMI_SUCCESS</code>	is returned upon successful call.
<code>None-zero</code>	is returned upon failure.

5.2 Energy Monitor (RAPL MSR)

Functions

- `esmi_status_t esmi_core_energy_get` (uint32_t core_ind, uint64_t *penergy)
Get the core energy for a given core.
- `esmi_status_t esmi_socket_energy_get` (uint32_t socket_idx, uint64_t *penergy)
Get the socket energy for a given socket.
- `esmi_status_t esmi_all_energies_get` (uint64_t *penergy)
Get energies of all cores in the system.

5.2.1 Detailed Description

Below functions provide interfaces to get the core energy value for a given core and to get the socket energy value for a given socket.

5.2.2 Function Documentation

5.2.2.1 `esmi_core_energy_get()`

```
esmi_status_t esmi_core_energy_get (
    uint32_t core_ind,
    uint64_t * penergy )
```

Get the core energy for a given core.

Given a core index `core_ind`, and a `penergy` argument for 64bit energy counter of that particular cpu, this function will read the energy counter of the given core and update the `penergy` in micro Joules.

Note: The energy status registers are accessed at core level. In a system with SMT enabled in BIOS, the sibling threads would report duplicate values. Aggregating the energy counters of the sibling threads is incorrect.

Parameters

in	<code>core_ind</code>	is a core index
in, out	<code>penergy</code>	Input buffer to return the core energy.

Return values

<code>ESMI_SUCCESS</code>	is returned upon successful call.
<code>None-zero</code>	is returned upon failure.

5.2.2.2 esmi_socket_energy_get()

```
esmi_status_t esmi_socket_energy_get (
    uint32_t socket_idx,
    uint64_t * penergy )
```

Get the socket energy for a given socket.

Given a socket index `socket_idx`, and a `penergy` argument for 64bit energy counter of a particular socket.

Updates the `penergy` with socket energy in micro Joules.

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>penergy</i>	Input buffer to return the socket energy.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.2.2.3 esmi_all_energies_get()

```
esmi_status_t esmi_all_energies_get (
    uint64_t * penergy )
```

Get energies of all cores in the system.

Given an argument for energy profile `penergy`, This function will read all core energies in an array `penergy` in micro Joules.

Parameters

in, out	<i>penergy</i>	Input buffer to return the energies of all cores. <code>penergy</code> should be allocated by user as below <code>(esmi_number_of_cpus_get()/esmi_threads_per_core_get()) * sizeof(uint64_t)</code>
---------	----------------	---

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3 HSMP System Statistics

Functions

- [esmi_status_t esmi_smu_fw_version_get](#) (struct [smu_fw_version](#) *smu_fw)
Get the SMU Firmware Version.
- [esmi_status_t esmi_prochot_status_get](#) (uint32_t socket_idx, uint32_t *prochot)
Get normalized status of the processor's PROCHOT status. 1 - PROCHOT active, 0 - PROCHOT inactive.
- [esmi_status_t esmi_fclk_mclk_get](#) (uint32_t socket_idx, uint32_t *fclk, uint32_t *mclk)
Get the Data Fabric clock and Memory clock in MHz, for a given socket index.
- [esmi_status_t esmi_cclk_limit_get](#) (uint32_t socket_idx, uint32_t *cclk)
Get the core clock (MHz) allowed by the most restrictive infrastructure limit at the time of the message.
- [esmi_status_t esmi_hsmp_proto_ver_get](#) (uint32_t *proto_ver)
Get the HSMP interface (protocol) version.
- [esmi_status_t esmi_socket_current_active_freq_limit_get](#) (uint32_t sock_ind, uint16_t *freq, char **src_type)
Get the current active frequency limit of the socket.
- [esmi_status_t esmi_socket_freq_range_get](#) (uint8_t sock_ind, uint16_t *fmax, uint16_t *fmin)
Get the Socket frequency range.
- [esmi_status_t esmi_current_freq_limit_core_get](#) (uint32_t core_id, uint32_t *freq)
Get the current active frequency limit of the core.

5.3.1 Detailed Description

Below functions to get HSMP System Statistics.

5.3.2 Function Documentation

5.3.2.1 esmi_smu_fw_version_get()

```
esmi_status_t esmi_smu_fw_version_get (
    struct smu\_fw\_version * smu_fw )
```

Get the SMU Firmware Version.

This function will return the SMU FW version at `smu_fw` Supported on all hsmp protocol versions

Parameters

in, out	smu_fw	Input buffer to return the smu firmware version.
---------	------------------------	--

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.2 esmi_prochot_status_get()

```
esmi_status_t esmi_prochot_status_get (
    uint32_t socket_idx,
    uint32_t * prochot )
```

Get normalized status of the processor's PROCHOT status. 1 - PROCHOT active, 0 - PROCHOT inactive.

Given a socket index `socket_idx` and this function will get PROCHOT at `prochot`. Supported on all hsmp protocol versions

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>prochot</i>	Input buffer to return the PROCHOT status.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.3 esmi_fclk_mclk_get()

```
esmi_status_t esmi_fclk_mclk_get (
    uint32_t socket_idx,
    uint32_t * fclk,
    uint32_t * mclk )
```

Get the Data Fabric clock and Memory clock in MHz, for a given socket index.

Given a socket index `socket_idx` and a pointer to a `uint32_t fclk` and `mclk`, this function will get the data fabric clock and memory clock. Supported on all hsmp protocol versions

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>fclk</i>	Input buffer to return the data fabric clock.
in, out	<i>mclk</i>	Input buffer to return the memory clock.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.4 esmi_cclk_limit_get()

```
esmi_status_t esmi_cclk_limit_get (
    uint32_t socket_idx,
    uint32_t * cclk )
```

Get the core clock (MHz) allowed by the most restrictive infrastructure limit at the time of the message.

Given a socket index `socket_idx` and a pointer to a `uint32_t cclk`, this function will get the core clock throttle limit. Supported on all hsmv protocol versions

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>cclk</i>	Input buffer to return the core clock throttle limit.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.5 esmi_hsmv_proto_ver_get()

```
esmi_status_t esmi_hsmv_proto_ver_get (
    uint32_t * proto_ver )
```

Get the HSMP interface (protocol) version.

This function will get the HSMP interface version at `proto_ver` Supported on all hsmv protocol versions

Parameters

in, out	<i>proto_ver</i>	Input buffer to return the hsmv protocol version.
---------	------------------	---

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.6 esmi_socket_current_active_freq_limit_get()

```
esmi_status_t esmi_socket_current_active_freq_limit_get (
    uint32_t sock_ind,
```

```
uint16_t * freq,
char ** src_type )
```

Get the current active frequency limit of the socket.

This function will get the socket frequency and source of this limit Supported on all hsmg protocol versions

Parameters

in	<i>sock_ind</i>	A socket index.
in, out	<i>freq</i>	Input buffer to return the frequency(MHz).
in, out	<i>src_type</i>	Input buffer to return the source of this limit

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.7 esmi_socket_freq_range_get()

```
esmi_status_t esmi_socket_freq_range_get (
    uint8_t sock_ind,
    uint16_t * fmax,
    uint16_t * fmin )
```

Get the Socket frequency range.

This function returns the socket frequency range, fmax and fmin. Supported only on hsmg protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index.
in, out	<i>fmax</i>	Input buffer to return the maximum frequency(MHz).
in, out	<i>fmin</i>	Input buffer to return the minimum frequency(MHz).

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.3.2.8 esmi_current_freq_limit_core_get()

```
esmi_status_t esmi_current_freq_limit_core_get (
    uint32_t core_id,
    uint32_t * freq )
```

Get the current active frequency limit of the core.

This function returns the core frequency limit for the specified core. Supported only on hsmv protocol version-5

Parameters

in	<i>core↔ _id</i>	Core index.
in, out	<i>freq</i>	Input buffer to return the core frequency limit(MHz)

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.4 Power Monitor

Functions

- [esmi_status_t esmi_socket_power_get](#) (uint32_t socket_idx, uint32_t *ppower)
Get the instantaneous power consumption of the provided socket.
- [esmi_status_t esmi_socket_power_cap_get](#) (uint32_t socket_idx, uint32_t *pcap)
Get the current power cap value for a given socket.
- [esmi_status_t esmi_socket_power_cap_max_get](#) (uint32_t socket_idx, uint32_t *pmax)
Get the maximum power cap value for a given socket.
- [esmi_status_t esmi_pwr_svi_telemetry_all_rails_get](#) (uint32_t sock_ind, uint32_t *power)
Get the SVI based power telemetry for all rails.

5.4.1 Detailed Description

Below functions provide interfaces to get the current power usage and Power Limits for a given socket.

5.4.2 Function Documentation

5.4.2.1 esmi_socket_power_get()

```
esmi_status_t esmi_socket_power_get (
    uint32_t socket_idx,
    uint32_t * ppower )
```

Get the instantaneous power consumption of the provided socket.

Given a socket index `socket_idx` and a pointer to a `uint32_t ppower`, this function will get the current power consumption (in milliwatts) to the `uint32_t` pointed to by `ppower`. Supported on all hsmg protocol versions

Parameters

in	<code>socket_idx</code>	a socket index
in, out	<code>ppower</code>	Input buffer to return power consumption in the socket.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.4.2.2 esmi_socket_power_cap_get()

```
esmi_status_t esmi_socket_power_cap_get (
```

```
uint32_t socket_idx,
uint32_t * pcap )
```

Get the current power cap value for a given socket.

This function will return the valid power cap `pcap` for a given socket `socket_idx`, this value will be used by the system to limit the power usage. Supported on all hsmf protocol versions

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>pcap</i>	Input buffer to return power limit on the socket, in milliwatts.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.4.2.3 esmi_socket_power_cap_max_get()

```
esmi_status_t esmi_socket_power_cap_max_get (
    uint32_t socket_idx,
    uint32_t * pmax )
```

Get the maximum power cap value for a given socket.

This function will return the maximum possible valid power cap `pmax` from a `socket_idx`. Supported on all hsmf protocol versions

Parameters

in	<i>socket_idx</i>	a socket index
in, out	<i>pmax</i>	Input buffer to return maximum power limit on socket, in milliwatts.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.4.2.4 esmi_pwr_svi_telemetry_all_rails_get()

```
esmi_status_t esmi_pwr_svi_telemetry_all_rails_get (
    uint32_t sock_ind,
    uint32_t * power )
```


Get the SVI based power telemetry for all rails.

This function returns the SVI based power telemetry for all rails. Supported only on hsmv protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index.
in, out	<i>power</i>	Input buffer to return the power(mW).

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.5 Power Control

Functions

- [esmi_status_t esmi_socket_power_cap_set](#) (uint32_t socket_idx, uint32_t pcap)
Set the power cap value for a given socket.
- [esmi_status_t esmi_pwr_efficiency_mode_set](#) (uint8_t sock_ind, uint8_t mode)
Set the power efficiency profile policy.

5.5.1 Detailed Description

This function provides a way to control Power Limit.

5.5.2 Function Documentation

5.5.2.1 esmi_socket_power_cap_set()

```
esmi_status_t esmi_socket_power_cap_set (
    uint32_t socket_idx,
    uint32_t pcap )
```

Set the power cap value for a given socket.

This function will set the power cap to the provided value `pcap`. This cannot be more than the value returned by [esmi_socket_power_cap_max_get\(\)](#).

Note: The power limit specified will be clipped to the maximum cTDP range for the processor. There is a limit on the minimum power that the processor can operate at, no further socket power reduction occurs if the limit is set below that minimum and also there are independent registers through HSMP and APML whichever is the most constraining between the two is enforced. Supported on all hsmp protocol versions.

Parameters

in	<i>socket_idx</i>	a socket index
in	<i>pcap</i>	a uint32_t that indicates the desired power cap, in milliwatts

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.5.2.2 esmi_pwr_efficiency_mode_set()

```
esmi_status_t esmi_pwr_efficiency_mode_set (
    uint8_t sock_ind,
    uint8_t mode )
```

Set the power efficiency profile policy.

This function will set the power efficiency mode. Supported only on hsmv protocol version-5

Power efficiency modes are:

0 = High performance mode: This mode favours core performance. In this mode all df pstates are available and default df pstate and DLWM algorithms are active.

1 = Power efficient mode: This mode limits the boost frequency available to the cores and restricts the DF P-States. This mode also monitors the system load to dynamically adjust performance for maximum power efficiency.

2 = IO performance mode: This mode sets up data fabric to maximize IO performance. This can result in lower core performance to increase the IO throughput.

3 = Balanced Memory Performance Mode: This mode biases the memory subsystem and Infinity Fabric™ performance towards efficiency, by lowering the frequency of the fabric and the width of the xGMI links under light traffic conditions. Core behavior is unaffected. There may be a performance impact under lightly loaded conditions for memory-bound applications compared to the default high performance mode. With higher memory and fabric load, the system becomes similar in performance to the default high performance mode.

Parameters

in	<i>sock_ind</i>	A socket index.
in	<i>mode</i>	Power efficiency mode to be set.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.6 Performance (Boost limit) Monitor

Functions

- [esmi_status_t esmi_core_boostlimit_get](#) (uint32_t cpu_ind, uint32_t *pboostlimit)
Get the boostlimit value for a given core.
- [esmi_status_t esmi_socket_c0_residency_get](#) (uint32_t socket_idx, uint32_t *pc0_residency)
Get the c0_residency value for a given socket.

5.6.1 Detailed Description

This function provides the current boostlimit value for a given core.

5.6.2 Function Documentation

5.6.2.1 esmi_core_boostlimit_get()

```
esmi_status_t esmi_core_boostlimit_get (
    uint32_t cpu_ind,
    uint32_t * pboostlimit )
```

Get the boostlimit value for a given core.

This function provides the frequency currently enforced through [esmi_core_boostlimit_set\(\)](#) and [esmi_socket_boostlimit_set\(\)](#) APIs for a particular `cpu_ind`. Supported on all hsmp protocol versions. Please note: there are independent registers through HSMP and APML. This message provides boost limit associated with HSMP only.

Parameters

in	<i>cpu_ind</i>	a cpu index
in, out	<i>pboostlimit</i>	Input buffer to return the boostlimit.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.6.2.2 esmi_socket_c0_residency_get()

```
esmi_status_t esmi_socket_c0_residency_get (
    uint32_t socket_idx,
    uint32_t * pc0_residency )
```

Get the `c0_residency` value for a given socket.

This function will return the socket's current `c0_residency` `pc0_residency` for a particular `socket_idx` Supported on all hsmf protocol versions

Parameters

in	<i>socket_idx</i>	a socket index provided.
in, out	<i>pc0_residency</i>	Input buffer to return the <code>c0_residency</code> .

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.7 Performance (Boost limit) Control

Functions

- [esmi_status_t esmi_core_boostlimit_set](#) (uint32_t cpu_ind, uint32_t boostlimit)
Set the boostlimit value for a given core.
- [esmi_status_t esmi_socket_boostlimit_set](#) (uint32_t socket_idx, uint32_t boostlimit)
Set the boostlimit value for a given socket.

5.7.1 Detailed Description

Below functions provide ways to control Boost limit values.

5.7.2 Function Documentation

5.7.2.1 esmi_core_boostlimit_set()

```
esmi_status_t esmi_core_boostlimit_set (
    uint32_t cpu_ind,
    uint32_t boostlimit )
```

Set the boostlimit value for a given core.

This function will set the boostlimit to the provided value `boostlimit` for a given `cpu` `cpu_ind`.

Note: Even though set boost limit provides ability to limit frequency on a core basis, if all the cores of a CCX are not programmed for the same boost limit frequency, then the lower-frequency cores are limited to a frequency resolution that can be as low as 20% of the requested frequency. If the specified boost limit frequency of a core is not supported, then the processor selects the next lower supported frequency. For processor with SMT enabled, writes to different APIC ids that map to the same physical core overwrite the previous write to that core. There are independent registers through HSMP and APML whichever is the most constraining between the two is enforced. Supported on all hsmp protocol versions

Parameters

in	<i>cpu_ind</i>	a cpu index is a given core to set the boostlimit
in	<i>boostlimit</i>	a uint32_t that indicates the desired boostlimit value of a given core. The maximum accepted value is 65535MHz(UINT16_MAX).

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.7.2.2 esmi_socket_boostlimit_set()

```
esmi_status_t esmi_socket_boostlimit_set (
    uint32_t socket_idx,
    uint32_t boostlimit )
```

Set the boostlimit value for a given socket.

This function will set the boostlimit to the provided value `boostlimit` for a given socket `socket_idx`. There are independent registers through HSMP and APML whichever is the most constraining between the two is enforced. Supported on all hsm protocol versions

Parameters

in	<i>socket_idx</i>	a socket index to set boostlimit.
in	<i>boostlimit</i>	a uint32_t that indicates the desired boostlimit value of a particular socket. The maximum accepted value is 65535MHz(UINT16_MAX).

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.8 ddr_bandwidth Monitor

Functions

- `esmi_status_t esmi_ddr_bw_get (struct ddr_bw_metrics *ddr_bw)`

Get the Theoretical maximum DDR Bandwidth in GB/s, Current utilized DDR Bandwidth in GB/s and Current utilized DDR Bandwidth as a percentage of theoretical maximum in a system. Supported only on hsmv protocol version >= 3.

5.8.1 Detailed Description

This function provides the DDR Bandwidth for a system

5.8.2 Function Documentation

5.8.2.1 esmi_ddr_bw_get()

```
esmi_status_t esmi_ddr_bw_get (
    struct ddr_bw_metrics * ddr_bw )
```

Get the Theoretical maximum DDR Bandwidth in GB/s, Current utilized DDR Bandwidth in GB/s and Current utilized DDR Bandwidth as a percentage of theoretical maximum in a system. Supported only on hsmv protocol version >= 3.

This function will return the DDR Bandwidth metrics `ddr_bw`

Parameters

<code>in, out</code>	<code>ddr_bw</code>	Input buffer to return the DDR bandwidth metrics, contains <code>max_bw</code> , <code>utilized_bw</code> and <code>utilized_pct</code> .
----------------------	---------------------	---

Return values

<code>ESMI_SUCCESS</code>	is returned upon successful call.
<code>None-zero</code>	is returned upon failure.

5.9 Temperature Query

Functions

- [esmi_status_t esmi_socket_temperature_get](#) (uint32_t sock_ind, uint32_t *ptmon)
Get temperature monitor for a given socket.

5.9.1 Detailed Description

This function provides the current temperature value in degree C.

5.9.2 Function Documentation

5.9.2.1 esmi_socket_temperature_get()

```
esmi_status_t esmi_socket_temperature_get (
    uint32_t sock_ind,
    uint32_t * ptmon )
```

Get temperature monitor for a given socket.

This function will return the socket's current temperature in milli degree celsius `ptmon` for a particular `sock_ind`. Supported only on hsmv protocol version-4

Parameters

in	<i>sock_ind</i>	a socket index provided.
in, out	<i>ptmon</i>	pointer to a uint32_t that indicates the possible tmon value.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.10 Dimm statistics

Functions

- [esmi_status_t esmi_dimm_temp_range_and_refresh_rate_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [temp_range_refresh_rate](#) *rate)
Get dimm temperature range and refresh rate.
- [esmi_status_t esmi_dimm_power_consumption_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [dimm_power](#) *dimm_pow)
Get dimm power consumption and update rate.
- [esmi_status_t esmi_dimm_thermal_sensor_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [dimm_thermal](#) *dimm_temp)
Get dimm thermal sensor.

5.10.1 Detailed Description

This function provides the dimm temperature, power and update rates.

5.10.2 Function Documentation

5.10.2.1 esmi_dimm_temp_range_and_refresh_rate_get()

```
esmi_status_t esmi_dimm_temp_range_and_refresh_rate_get (
    uint8_t sock_ind,
    uint8_t dimm_addr,
    struct temp_range_refresh_rate * rate )
```

Get dimm temperature range and refresh rate.

This function returns the per DIMM temperature range and refresh rate from the MR4 register. Supported only on hsmp protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index through which the DIMM can be accessed
in	<i>dimm_addr</i>	DIMM identifier, follow "HSMP DIMM Address encoding".
in, out	<i>rate</i>	Input buffer of type struct temp_range_refresh_rate with refresh rate and temp range.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.10.2.2 esmi_dimm_power_consumption_get()

```
esmi_status_t esmi_dimm_power_consumption_get (
    uint8_t sock_ind,
    uint8_t dimm_addr,
    struct dimm_power * dimm_pow )
```

Get dimm power consumption and update rate.

This function returns the DIMM power and update rate Supported only on hsmp protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index through which the DIMM can be accessed.
in	<i>dimm_addr</i>	DIMM identifier, follow "HSMP DIMM Address encoding".
in, out	<i>dimm_pow</i>	Input buffer of type struct dimm_power containing power(mW), update rate(ms) and dimm address.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.10.2.3 esmi_dimm_thermal_sensor_get()

```
esmi_status_t esmi_dimm_thermal_sensor_get (
    uint8_t sock_ind,
    uint8_t dimm_addr,
    struct dimm_thermal * dimm_temp )
```

Get dimm thermal sensor.

This function will return the DIMM thermal sensor(2 sensors per DIMM) and update rate Supported only on hsmp protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index through which the DIMM can be accessed.
in	<i>dimm_addr</i>	DIMM identifier, follow "HSMP DIMM Address encoding".
in, out	<i>dimm_temp</i>	Input buffer of type struct dimm_thermal which contains temperature(°C), update rate(ms) and dimm address Update rate value can vary from 0 to 511ms. Update rate of "0" means last update was < 1ms and 511ms means update was >= 511ms.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.11 xGMI bandwidth control

Functions

- [esmi_status_t esmi_xgmi_width_set](#) (uint8_t min, uint8_t max)
Set xgmi width for a multi socket system. values range from 0 to 2.

5.11.1 Detailed Description

This function provides a way to control width of the xgmi links connected in multisocket systems.

5.11.2 Function Documentation

5.11.2.1 esmi_xgmi_width_set()

```
esmi_status_t esmi_xgmi_width_set (
    uint8_t min,
    uint8_t max )
```

Set xgmi width for a multi socket system. values range from 0 to 2.

0 => 4 lanes on family 19h model 10h and 2 lanes on other models.

1 => 8 lanes.

2 => 16 lanes.

Supported on all hsmp protocol versions.

This function will set the xgmi width `min` and `max` for all the sockets in the system

Parameters

in	<i>min</i>	minimum xgmi link width, varies from 0 to 2 with min <= max.
in	<i>max</i>	maximum xgmi link width, varies from 0 to 2.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.12 GMI3 width control

Functions

- [esmi_status_t esmi_gmi3_link_width_range_set](#) (uint8_t sock_ind, uint8_t min_link_width, uint8_t max_↔ link_width)
Set gmi3 width.

5.12.1 Detailed Description

This function provides a way to control global memory interconnect link width.

5.12.2 Function Documentation

5.12.2.1 esmi_gmi3_link_width_range_set()

```
esmi_status_t esmi_gmi3_link_width_range_set (
    uint8_t sock_ind,
    uint8_t min_link_width,
    uint8_t max_link_width )
```

Set gmi3 width.

This function will set the global memory interconnect width. Values can be 0, 1 or 2.

0 => Quarter width

1 => Half width

2 => Full width

Supported only on hsmv protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index.
in	<i>min_link_width</i>	Minimum link width to be set.
in	<i>max_link_width</i>	Maximum link width to be set.

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13 APB and LCLK level control

Functions

- [esmi_status_t esmi_apb_enable](#) (uint32_t sock_ind)
Enable automatic P-state selection.
- [esmi_status_t esmi_apb_disable](#) (uint32_t sock_ind, uint8_t pstate)
Set data fabric P-state to user specified value.
- [esmi_status_t esmi_socket_lclk_dpm_level_set](#) (uint32_t sock_ind, uint8_t nbio_id, uint8_t min, uint8_t max)
Set lclk dpm level.
- [esmi_status_t esmi_socket_lclk_dpm_level_get](#) (uint8_t sock_ind, uint8_t nbio_id, struct [dpm_level](#) *nbio)
Get lclk dpm level.
- [esmi_status_t esmi_pcie_link_rate_set](#) (uint8_t sock_ind, uint8_t rate_ctrl, uint8_t *prev_mode)
Set pcie link rate.
- [esmi_status_t esmi_df_pstate_range_set](#) (uint8_t sock_ind, uint8_t max_pstate, uint8_t min_pstate)
Set data fabric pstate range.

5.13.1 Detailed Description

This functions provides a way to control APB and lclk values.

5.13.2 Function Documentation

5.13.2.1 esmi_apb_enable()

```
esmi_status_t esmi_apb_enable (
    uint32_t sock_ind )
```

Enable automatic P-state selection.

Given a socket index `sock_ind`, this function will enable performance boost algorithm. By default, an algorithm adjusts DF P-States automatically in order to optimize performance. However, this default may be changed to a fixed DF P-State through a CBS option at boottime. APBDisable may also be used to disable this algorithm and force a fixed DF P-State. Supported on all hsmpt protocol versions

NOTE: While the socket is in PC6 or if PROCHOT_L is asserted, the lowest DF P-State (highest value) is enforced regardless of the APBEnable/APBDisable state.

Parameters

in	<code>sock_ind</code>	a socket index
----	-----------------------	----------------

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13.2.2 esmi_apb_disable()

```
esmi_status_t esmi_apb_disable (
    uint32_t sock_ind,
    uint8_t pstate )
```

Set data fabric P-state to user specified value.

This function will set the desired P-state at `pstate`. Acceptable values for the P-state are 0(highest) - 2 (lowest) If the PC6 or PROCHOT_L is asserted, then the lowest DF pstate is enforced regardless of the APBenable/APBdiable states. Supported on all hsmp protocol versions.

Parameters

in	<i>sock_ind</i>	a socket index
in	<i>pstate</i>	a uint8_t that indicates the desired P-state to set.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13.2.3 esmi_socket_lclk_dpm_level_set()

```
esmi_status_t esmi_socket_lclk_dpm_level_set (
    uint32_t sock_ind,
    uint8_t nbio_id,
    uint8_t min,
    uint8_t max )
```

Set lclk dpm level.

This function will set the lclk dpm level / nbio pstate for the specified `nbio_id` in a specified socket `sock_ind` with provided values `min` and `max`. Supported on hsmp protocol version ≥ 2

Parameters

in	<i>sock_ind</i>	socket index.
in	<i>nbio_id</i>	northbridge number varies from 0 to 3.
in	<i>min</i>	pstate minimum value, varies from 0(lowest) to 3(highest) with $\text{min} \leq \text{max}$
in	<i>max</i>	pstate maximum value, varies from 0 to 3.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13.2.4 esmi_socket_lclk_dpm_level_get()

```
esmi_status_t esmi_socket_lclk_dpm_level_get (
    uint8_t sock_ind,
    uint8_t nbio_id,
    struct dpm_level * nbio )
```

Get lclk dpm level.

This function will get the lclk dpm level. DPM level is an encoding to represent PCIe link frequency. DPM levels can be set from APML also. This API gives current levels which may have been set from either APML or HSMP.

Supported in hsmg protocol version-5.

Parameters

in	<i>sock_ind</i>	Socket index
in	<i>nbio_id</i>	NBIO id(0-3)
in, out	<i>nbio</i>	Input buffer of struct dpm_level type to hold min and max dpm levels

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13.2.5 esmi_pcie_link_rate_set()

```
esmi_status_t esmi_pcie_link_rate_set (
    uint8_t sock_ind,
    uint8_t rate_ctrl,
    uint8_t * prev_mode )
```

Set pcie link rate.

This function will set the pcie link rate to gen4/5 or auto detection based on bandwidth utilisation. Values are: 0 => auto detect bandwidth utilisation and set link rate

1 => Limit at gen4 rate

2 => Limit at gen5 rate

Supported only on hsmg protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index.
in	<i>rate_ctrl</i>	Control value to be set.
in, out	<i>prev_mode</i>	Input buffer to hold the previous mode.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.13.2.6 esmi_df_pstate_range_set()

```
esmi_status_t esmi_df_pstate_range_set (
    uint8_t sock_ind,
    uint8_t max_pstate,
    uint8_t min_pstate )
```

Set data fabric pstate range.

This function will set the max and min pstates for the data fabric. Acceptable values for the P-state are 0(highest) - 2 (lowest) with max <= min. DF pstate range can be set from both HSMP and APML, the most recent of the two is enforced. Supported only on hsmp protocol version-5

Parameters

in	<i>sock_ind</i>	a socket index.
in	<i>max_pstate</i>	Maximum pstate value to be set.
in	<i>min_pstate</i>	Minimum pstate value to be set.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.14 Bandwidth Monitor

Functions

- `esmi_status_t esmi_current_io_bandwidth_get` (uint8_t sock_ind, struct link_id_bw_type link, uint32_t *io_bw)
Get IO bandwidth on IO link.
- `esmi_status_t esmi_current_xgmi_bw_get` (struct link_id_bw_type link, uint32_t *xgmi_bw)
Get xGMI bandwidth.

5.14.1 Detailed Description

This function provides the IO and xGMI bandwidth.

5.14.2 Function Documentation

5.14.2.1 esmi_current_io_bandwidth_get()

```
esmi_status_t esmi_current_io_bandwidth_get (
    uint8_t sock_ind,
    struct link_id_bw_type link,
    uint32_t * io_bw )
```

Get IO bandwidth on IO link.

This function returns the IO Aggregate bandwidth for the given link id. Supported only on hsmp protocol version-5

Parameters

in	<i>sock_ind</i>	Socket index.
in	<i>link</i>	structure containing link_id(Link encoding values of given link) and bwtype info.
in, out	<i>io_bw</i>	Input buffer for bandwidth data in Mbps.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.14.2.2 esmi_current_xgmi_bw_get()

```
esmi_status_t esmi_current_xgmi_bw_get (
    struct link_id_bw_type link,
    uint32_t * xgmi_bw )
```

Get xGMI bandwidth.

This function will read xGMI bandwidth in Mbps for the specified link and bandwidth type in a multi socket system. Supported only on hsmv protocol version-5

Parameters

in	<i>link</i>	structure containing link_id(Link encoding values of given link) and bwtype info.
in, out	<i>xgmi_bw</i>	Input buffer for bandwidth data in Mbps.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.15 Metrics Table

Functions

- [esmi_status_t esmi_metrics_table_version_get](#) (uint32_t *metrics_version)
Get metrics table version.
- [esmi_status_t esmi_metrics_table_get](#) (uint8_t sock_ind, struct hsmp_metric_table *metrics_table)
Get metrics table.
- [esmi_status_t esmi_dram_address_metrics_table_get](#) (uint8_t sock_ind, uint64_t *dram_addr)
Get the DRAM address for the metrics table.

5.15.1 Detailed Description

The following functions are assigned for CPU relative functionality, which is expected to be compatible with Epyc products.

5.15.2 Function Documentation

5.15.2.1 esmi_metrics_table_version_get()

```
esmi_status_t esmi_metrics_table_version_get (
    uint32_t * metrics_version )
```

Get metrics table version.

Get the version number[31:0] of metrics table

Parameters

in, out	<i>metrics_version</i>	input buffer to return the metrics table version.
---------	------------------------	---

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>Non-zero</i>	is returned upon failure.

5.15.2.2 esmi_metrics_table_get()

```
esmi_status_t esmi_metrics_table_get (
    uint8_t sock_ind,
    struct hsmp_metric_table * metrics_table )
```

Get metrics table.

Read the metrics table

Parameters

in	<i>sock_ind</i>	Socket index.
in, out	<i>metrics_table</i>	input buffer to return the metrics table.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>Non-zero</i>	is returned upon failure.

5.15.2.3 esmi_dram_address_metrics_table_get()

```
esmi_status_t esmi_dram_address_metrics_table_get (
    uint8_t sock_ind,
    uint64_t * dram_addr )
```

Get the DRAM address for the metrics table.

Get DRAM address for Metric table transfer

Parameters

in	<i>sock_ind</i>	Socket index.
in, out	<i>dram_addr</i>	64-bit DRAM address

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>Non-zero</i>	is returned upon failure.

5.16 Test HSMP mailbox

Functions

- [esmi_status_t esmi_test_hsmp_mailbox](#) (uint8_t sock_ind, uint32_t *data)
Test HSMP mailbox interface.

5.16.1 Detailed Description

This is used to check if the HSMP interface is functioning correctly. Increments the input argument value by 1.

5.16.2 Function Documentation

5.16.2.1 esmi_test_hsmp_mailbox()

```
esmi_status_t esmi_test_hsmp_mailbox (
    uint8_t sock_ind,
    uint32_t * data )
```

Test HSMP mailbox interface.

[31:0] = input value

Parameters

in	<i>sock_ind</i>	:Socket index.
	<i>[in/out]</i>	data : input buffer to send input value and to get the output value

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>Non-zero</i>	is returned upon failure.

5.17 Auxiliary functions

Functions

- [esmi_status_t esmi_cpu_family_get](#) (uint32_t *family)
Get the CPU family.
- [esmi_status_t esmi_cpu_model_get](#) (uint32_t *model)
Get the CPU model.
- [esmi_status_t esmi_threads_per_core_get](#) (uint32_t *threads)
Get the number of threads per core in the system.
- [esmi_status_t esmi_number_of_cpus_get](#) (uint32_t *cpus)
Get the number of cpus available in the system.
- [esmi_status_t esmi_number_of_sockets_get](#) (uint32_t *sockets)
Get the total number of sockets available in the system.
- [esmi_status_t esmi_first_online_core_on_socket](#) (uint32_t socket_idx, uint32_t *pcore_ind)
Get the first online core on a given socket.
- char * [esmi_get_err_msg](#) (esmi_status_t esmi_err)
Get the error string message for esmi errors.

5.17.1 Detailed Description

Below functions provide interfaces to get the total number of cores and sockets available and also to get the first online core on a given socket in the system.

5.17.2 Function Documentation

5.17.2.1 esmi_cpu_family_get()

```
esmi_status_t esmi_cpu_family_get (
    uint32_t * family )
```

Get the CPU family.

Parameters

<i>in, out</i>	<i>family</i>	Input buffer to return the cpu family.
----------------	---------------	--

Return values

ESMI_SUCCESS	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.2 esmi_cpu_model_get()

```
esmi_status_t esmi_cpu_model_get (
    uint32_t * model )
```

Get the CPU model.

Parameters

<i>in, out</i>	<i>model</i>	Input buffer to return the cpu model.
----------------	--------------	---------------------------------------

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.3 esmi_threads_per_core_get()

```
esmi_status_t esmi_threads_per_core_get (
    uint32_t * threads )
```

Get the number of threads per core in the system.

Parameters

<i>in, out</i>	<i>threads</i>	input buffer to return number of SMT threads.
----------------	----------------	---

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.4 esmi_number_of_cpus_get()

```
esmi_status_t esmi_number_of_cpus_get (
    uint32_t * cpus )
```

Get the number of cpus available in the system.

Parameters

<i>in, out</i>	<i>cpus</i>	input buffer to return number of cpus, reported by nproc (including threads in case of SMT enable).
----------------	-------------	---

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.5 esmi_number_of_sockets_get()

```
esmi_status_t esmi_number_of_sockets_get (
    uint32_t * sockets )
```

Get the total number of sockets available in the system.

Parameters

in, out	<i>sockets</i>	input buffer to return number of sockets.
---------	----------------	---

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.6 esmi_first_online_core_on_socket()

```
esmi_status_t esmi_first_online_core_on_socket (
    uint32_t socket_idx,
    uint32_t * pcore_ind )
```

Get the first online core on a given socket.

Parameters

in	<i>socket_idx</i>	a socket index provided.
in, out	<i>pcore_ind</i>	input buffer to return the index of first online core in the socket.

Return values

<i>ESMI_SUCCESS</i>	is returned upon successful call.
<i>None-zero</i>	is returned upon failure.

5.17.2.7 esmi_get_err_msg()

```
char* esmi_get_err_msg (
    esmi_status_t esmi_err )
```

Get the error string message for esmi errors.

Get the error message for the esmi error numbers

Parameters

in	<i>esmi_err</i>	is a esmi error number
----	-----------------	------------------------

Return values

<i>char*</i>	value returned upon successful call.
--------------	--------------------------------------

Chapter 6

Data Structure Documentation

6.1 ddr_bw_metrics Struct Reference

DDR bandwidth metrics.

```
#include <e_smi.h>
```

Data Fields

- uint32_t [max_bw](#)
DDR Maximum theoretical bandwidth in GB/s.
- uint32_t [utilized_bw](#)
DDR bandwidth utilization in GB/s.
- uint32_t [utilized_pct](#)
DDR bandwidth utilization in % of theoretical max.

6.1.1 Detailed Description

DDR bandwidth metrics.

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.2 dimm_power Struct Reference

DIMM Power(mW), power update rate(ms) and dimm address.

```
#include <e_smi.h>
```

Data Fields

- uint16_t [power](#): 15
Dimm power consumption[31:17](15 bits data)
- uint16_t [update_rate](#): 9
Time since last update[16:8](9 bit data)
- uint8_t [dimm_addr](#)
Dimm address[7:0](8 bit data)

6.2.1 Detailed Description

DIMM Power(mW), power update rate(ms) and dimm address.

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.3 dimm_thermal Struct Reference

DIMM temperature(°C) and update rate(ms) and dimm address.

```
#include <e_smi.h>
```

Data Fields

- uint16_t [sensor](#): 11
Dimm thermal sensor[31:21](11 bit data)
- uint16_t [update_rate](#): 9
Time since last update[16:8](9 bit data)
- uint8_t [dimm_addr](#)
Dimm address[7:0](8 bit data)
- float [temp](#)
temperature in degree celcius

6.3.1 Detailed Description

DIMM temperature(°C) and update rate(ms) and dimm address.

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.4 dpm_level Struct Reference

max and min LCLK DPM level on a given NBIO ID. Valid max and min DPM level values are 0 - 1.

```
#include <e_smi.h>
```

Data Fields

- [uint8_t max_dpm_level](#)
Max LCLK DPM level[15:8](8 bit data)
- [uint8_t min_dpm_level](#)
Min LCLK DPM level[7:0](8 bit data)

6.4.1 Detailed Description

max and min LCLK DPM level on a given NBIO ID. Valid max and min DPM level values are 0 - 1.

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.5 link_id_bw_type Struct Reference

LINK name and Bandwidth type Information. It contains link names i.e valid link names are "P0", "P1", "P2", "P3", "P4", "G0", "G1", "G2", "G3", "G4", "G5", "G6", "G7" Valid bandwidth types 1 (Aggregate_BW), 2 (Read BW), 4 (Write BW).

```
#include <e_smi.h>
```

Data Fields

- [io_bw_encoding bw_type](#)
Bandwidth Type Information [1, 2, 4].
- [char * link_name](#)
Link name [P0, P1, G0, G1 etc].

6.5.1 Detailed Description

LINK name and Bandwidth type Information. It contains link names i.e valid link names are "P0", "P1", "P2", "P3", "P4", "G0", "G1", "G2", "G3", "G4", "G5", "G6", "G7" Valid bandwidth types 1 (Aggregate_BW), 2 (Read BW), 4 (Write BW).

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.6 smu_fw_version Struct Reference

Deconstruct raw uint32_t into SMU firmware major and minor version numbers.

```
#include <e_smi.h>
```

Data Fields

- `uint8_t debug`
SMU fw Debug version number.
- `uint8_t minor`
SMU fw Minor version number.
- `uint8_t major`
SMU fw Major version number.
- `uint8_t unused`
reserved fields

6.6.1 Detailed Description

Deconstruct raw `uint32_t` into SMU firmware major and minor version numbers.

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

6.7 temp_range_refresh_rate Struct Reference

temperature range and refresh rate metrics of a DIMM

```
#include <e_smi.h>
```

Data Fields

- `uint8_t range`: 3
temp range[2:0](3 bit data)
- `uint8_t ref_rate`: 1
DDR refresh rate mode[3](1 bit data)

6.7.1 Detailed Description

temperature range and refresh rate metrics of a DIMM

The documentation for this struct was generated from the following file:

- [e_smi.h](#)

Chapter 7

File Documentation

7.1 e_smi.h File Reference

```
#include <stdbool.h>
#include <asm/amd_hsmpt.h>
```

Data Structures

- struct [smu_fw_version](#)
Deconstruct raw uint32_t into SMU firmware major and minor version numbers.
- struct [ddr_bw_metrics](#)
DDR bandwidth metrics.
- struct [temp_range_refresh_rate](#)
temperature range and refresh rate metrics of a DIMM
- struct [dimm_power](#)
DIMM Power(mW), power update rate(ms) and dimm address.
- struct [dimm_thermal](#)
DIMM temperature(°C) and update rate(ms) and dimm address.
- struct [link_id_bw_type](#)
LINK name and Bandwidth type Information.It contains link names i.e valid link names are "P0", "P1", "P2", "P3", "P4", "G0", "G1", "G2", "G3", "G4", "G5", "G6", "G7" Valid bandwidth types 1(Aggregate_BW), 2 (Read BW), 4 (Write BW).
- struct [dpm_level](#)
max and min LCLK DPM level on a given NBIO ID. Valid max and min DPM level values are 0 - 1.

Macros

- #define [ENERGY_DEV_NAME](#) "amd_energy"
Supported Energy driver name.
- #define [HSMP_CHAR_DEVFILE_NAME](#) "/dev/hsmpt"
HSMP device path.
- #define [HSMP_METRICTABLE_PATH](#) "/sys/devices/platform/amd_hsmpt"
HSMP MetricTable sysfs path.
- #define [ARRAY_SIZE](#)(arr) (sizeof(arr) / sizeof((arr)[0]))
macro to calculate size
- #define [BIT](#)(N) (1 << N)
macro for mask

Enumerations

- enum `io_bw_encoding` { `AGG_BW` = BIT(0), `RD_BW` = BIT(1), `WR_BW` = BIT(2) }
xGMI Bandwidth Encoding types
- enum `esmi_status_t` {
`ESMI_SUCCESS` = 0, `ESMI_INITIALIZED` = 0, `ESMI_NO_ENERGY_DRV`, `ESMI_NO_MSR_DRV`,
`ESMI_NO_HSMP_DRV`, `ESMI_NO_HSMP_SUP`, `ESMI_NO_DRV`, `ESMI_FILE_NOT_FOUND`,
`ESMI_DEV_BUSY`, `ESMI_PERMISSION`, `ESMI_NOT_SUPPORTED`, `ESMI_FILE_ERROR`,
`ESMI_INTERRUPTED`, `ESMI_IO_ERROR`, `ESMI_UNEXPECTED_SIZE`, `ESMI_UNKNOWN_ERROR`,
`ESMI_ARG_PTR_NULL`, `ESMI_NO_MEMORY`, `ESMI_NOT_INITIALIZED`, `ESMI_INVALID_INPUT`,
`ESMI_HSMP_TIMEOUT`, `ESMI_NO_HSMP_MSG_SUP` }
Error codes returned by E-SMI functions.

Functions

- `esmi_status_t esmi_init` (void)
Initialize the library, validates the dependencies.
- void `esmi_exit` (void)
Clean up any allocation done during init.
- `esmi_status_t esmi_core_energy_get` (uint32_t core_ind, uint64_t *penergy)
Get the core energy for a given core.
- `esmi_status_t esmi_socket_energy_get` (uint32_t socket_idx, uint64_t *penergy)
Get the socket energy for a given socket.
- `esmi_status_t esmi_all_energies_get` (uint64_t *penergy)
Get energies of all cores in the system.
- `esmi_status_t esmi_smu_fw_version_get` (struct `smu_fw_version` *smu_fw)
Get the SMU Firmware Version.
- `esmi_status_t esmi_prochot_status_get` (uint32_t socket_idx, uint32_t *prochot)
Get normalized status of the processor's PROCHOT status. 1 - PROCHOT active, 0 - PROCHOT inactive.
- `esmi_status_t esmi_fclk_mclk_get` (uint32_t socket_idx, uint32_t *fclk, uint32_t *mclk)
Get the Data Fabric clock and Memory clock in MHz, for a given socket index.
- `esmi_status_t esmi_cclk_limit_get` (uint32_t socket_idx, uint32_t *cclock)
Get the core clock (MHz) allowed by the most restrictive infrastructure limit at the time of the message.
- `esmi_status_t esmi_hsmp_proto_ver_get` (uint32_t *proto_ver)
Get the HSMP interface (protocol) version.
- `esmi_status_t esmi_socket_current_active_freq_limit_get` (uint32_t sock_ind, uint16_t *freq, char **src_↵ type)
Get the current active frequency limit of the socket.
- `esmi_status_t esmi_socket_freq_range_get` (uint8_t sock_ind, uint16_t *fmax, uint16_t *fmin)
Get the Socket frequency range.
- `esmi_status_t esmi_current_freq_limit_core_get` (uint32_t core_id, uint32_t *freq)
Get the current active frequency limit of the core.
- `esmi_status_t esmi_socket_power_get` (uint32_t socket_idx, uint32_t *ppower)
Get the instantaneous power consumption of the provided socket.
- `esmi_status_t esmi_socket_power_cap_get` (uint32_t socket_idx, uint32_t *pcap)
Get the current power cap value for a given socket.
- `esmi_status_t esmi_socket_power_cap_max_get` (uint32_t socket_idx, uint32_t *pmax)
Get the maximum power cap value for a given socket.
- `esmi_status_t esmi_pwr_svi_telemetry_all_rails_get` (uint32_t sock_ind, uint32_t *power)
Get the SVI based power telemetry for all rails.
- `esmi_status_t esmi_socket_power_cap_set` (uint32_t socket_idx, uint32_t pcap)

- Set the power cap value for a given socket.*

 - [esmi_status_t esmi_pwr_efficiency_mode_set](#) (uint8_t sock_ind, uint8_t mode)
- Set the power efficiency profile policy.*

 - [esmi_status_t esmi_core_boostlimit_get](#) (uint32_t cpu_ind, uint32_t *pboostlimit)

Get the boostlimit value for a given core.

 - [esmi_status_t esmi_socket_c0_residency_get](#) (uint32_t socket_idx, uint32_t *pc0_residency)

Get the c0_residency value for a given socket.

 - [esmi_status_t esmi_core_boostlimit_set](#) (uint32_t cpu_ind, uint32_t boostlimit)

Set the boostlimit value for a given core.

 - [esmi_status_t esmi_socket_boostlimit_set](#) (uint32_t socket_idx, uint32_t boostlimit)

Set the boostlimit value for a given socket.

 - [esmi_status_t esmi_ddr_bw_get](#) (struct [ddr_bw_metrics](#) *ddr_bw)

Get the Theoretical maximum DDR Bandwidth in GB/s, Current utilized DDR Bandwidth in GB/s and Current utilized DDR Bandwidth as a percentage of theoretical maximum in a system. Supported only on hsmv protocol version >= 3.

 - [esmi_status_t esmi_socket_temperature_get](#) (uint32_t sock_ind, uint32_t *ptmon)

Get temperature monitor for a given socket.

 - [esmi_status_t esmi_dimm_temp_range_and_refresh_rate_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [temp_range_refresh_rate](#) *rate)

Get dimm temperature range and refresh rate.

 - [esmi_status_t esmi_dimm_power_consumption_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [dimm_power](#) *dimm_pow)

Get dimm power consumption and update rate.

 - [esmi_status_t esmi_dimm_thermal_sensor_get](#) (uint8_t sock_ind, uint8_t dimm_addr, struct [dimm_thermal](#) *dimm_temp)

Get dimm thermal sensor.

 - [esmi_status_t esmi_xgmi_width_set](#) (uint8_t min, uint8_t max)

Set xgmi width for a multi socket system. values range from 0 to 2.

 - [esmi_status_t esmi_gmi3_link_width_range_set](#) (uint8_t sock_ind, uint8_t min_link_width, uint8_t max_link_width)

Set gmi3 width.

 - [esmi_status_t esmi_apb_enable](#) (uint32_t sock_ind)

Enable automatic P-state selection.

 - [esmi_status_t esmi_apb_disable](#) (uint32_t sock_ind, uint8_t pstate)

Set data fabric P-state to user specified value.

 - [esmi_status_t esmi_socket_lclk_dpm_level_set](#) (uint32_t sock_ind, uint8_t nbio_id, uint8_t min, uint8_t max)

Set lclk dpm level.

 - [esmi_status_t esmi_socket_lclk_dpm_level_get](#) (uint8_t sock_ind, uint8_t nbio_id, struct [dpm_level](#) *nbio)

Get lclk dpm level.

 - [esmi_status_t esmi_pcie_link_rate_set](#) (uint8_t sock_ind, uint8_t rate_ctrl, uint8_t *prev_mode)

Set pcie link rate.

 - [esmi_status_t esmi_df_pstate_range_set](#) (uint8_t sock_ind, uint8_t max_pstate, uint8_t min_pstate)

Set data fabric pstate range.

 - [esmi_status_t esmi_current_io_bandwidth_get](#) (uint8_t sock_ind, struct [link_id_bw_type](#) link, uint32_t *io_bw)

Get IO bandwidth on IO link.

 - [esmi_status_t esmi_current_xgmi_bw_get](#) (struct [link_id_bw_type](#) link, uint32_t *xgmi_bw)

Get xGMI bandwidth.

 - [esmi_status_t esmi_metrics_table_version_get](#) (uint32_t *metrics_version)

Get metrics table version.

 - [esmi_status_t esmi_metrics_table_get](#) (uint8_t sock_ind, struct [hsmv_metric_table](#) *metrics_table)

Get metrics table.

- `esmi_status_t esmi_dram_address_metrics_table_get` (uint8_t sock_ind, uint64_t *dram_addr)
Get the DRAM address for the metrics table.
- `esmi_status_t esmi_test_hsmp_mailbox` (uint8_t sock_ind, uint32_t *data)
Test HSMP mailbox interface.
- `esmi_status_t esmi_cpu_family_get` (uint32_t *family)
Get the CPU family.
- `esmi_status_t esmi_cpu_model_get` (uint32_t *model)
Get the CPU model.
- `esmi_status_t esmi_threads_per_core_get` (uint32_t *threads)
Get the number of threads per core in the system.
- `esmi_status_t esmi_number_of_cpus_get` (uint32_t *cpus)
Get the number of cpus available in the system.
- `esmi_status_t esmi_number_of_sockets_get` (uint32_t *sockets)
Get the total number of sockets available in the system.
- `esmi_status_t esmi_first_online_core_on_socket` (uint32_t socket_idx, uint32_t *pcore_ind)
Get the first online core on a given socket.
- `char * esmi_get_err_msg` (esmi_status_t esmi_err)
Get the error string message for esmi errors.

7.1.1 Detailed Description

Main header file for the E-SMI library. All required function, structure, enum, etc. definitions should be defined in this file.

This header file contains the following: APIs prototype of the APIs exported by the E-SMI library. Description of the API, arguments and return values. The Error codes returned by the API.

7.1.2 Enumeration Type Documentation

7.1.2.1 io_bw_encoding

enum `io_bw_encoding`

xGMI Bandwidth Encoding types

Enumerator

AGG_BW	Aggregate Bandwidth.
RD_BW	Read Bandwidth.
WR_BW	Write Bandwdith.

7.1.2.2 esmi_status_t

enum `esmi_status_t`

Error codes returned by E-SMI functions.

Enumerator

ESMI_SUCCESS	Operation was successful.
ESMI_INITIALIZED	ESMI initialized successfully.
ESMI_NO_ENERGY_DRV	Energy driver not found.
ESMI_NO_MSR_DRV	MSR driver not found.
ESMI_NO_HSMP_DRV	HSMP driver not found.
ESMI_NO_HSMP_SUP	HSMP not supported.
ESMI_NO_DRV	No Energy and HSMP driver present.
ESMI_FILE_NOT_FOUND	file or directory not found
ESMI_DEV_BUSY	Device or resource busy.
ESMI_PERMISSION	Many functions require root access to run. Permission denied/EACCESS file error.
ESMI_NOT_SUPPORTED	The requested information or action is not available for the given input, on the given system
ESMI_FILE_ERROR	Problem accessing a file. This may because the operation is not supported by the Linux kernel version running on the executing machine
ESMI_INTERRUPTED	execution of function An interrupt occurred during
ESMI_IO_ERROR	An input or output error.
ESMI_UNEXPECTED_SIZE	was read An unexpected amount of data
ESMI_UNKNOWN_ERROR	An unknown error occurred.
ESMI_ARG_PTR_NULL	Parsed argument is invalid.
ESMI_NO_MEMORY	Not enough memory to allocate.
ESMI_NOT_INITIALIZED	ESMI path not initialized.
ESMI_INVALID_INPUT	Input value is invalid.
ESMI_HSMP_TIMEOUT	HSMP message is timedout.
ESMI_NO_HSMP_MSG_SUP	HSMP message/feature not supported.

Index

- APB and LCLK level control, [39](#)
 - [esmi_apb_disable](#), [40](#)
 - [esmi_apb_enable](#), [39](#)
 - [esmi_df_pstate_range_set](#), [42](#)
 - [esmi_pcie_link_rate_set](#), [41](#)
 - [esmi_socket_lclk_dpm_level_get](#), [41](#)
 - [esmi_socket_lclk_dpm_level_set](#), [40](#)
- Auxiliary functions, [49](#)
 - [esmi_cpu_family_get](#), [49](#)
 - [esmi_cpu_model_get](#), [49](#)
 - [esmi_first_online_core_on_socket](#), [51](#)
 - [esmi_get_err_msg](#), [51](#)
 - [esmi_number_of_cpus_get](#), [50](#)
 - [esmi_number_of_sockets_get](#), [51](#)
 - [esmi_threads_per_core_get](#), [50](#)
- Bandwidth Monitor, [43](#)
 - [esmi_current_io_bandwidth_get](#), [43](#)
 - [esmi_current_xgmi_bw_get](#), [43](#)
- ddr_bandwidth Monitor, [33](#)
 - [esmi_ddr_bw_get](#), [33](#)
- ddr_bw_metrics, [53](#)
- Dimm statistics, [35](#)
 - [esmi_dimm_power_consumption_get](#), [35](#)
 - [esmi_dimm_temp_range_and_refresh_rate_get](#), [35](#)
 - [esmi_dimm_thermal_sensor_get](#), [36](#)
- dimm_power, [53](#)
- dimm_thermal, [54](#)
- dpm_level, [54](#)
- e_smi.h, [57](#)
 - [esmi_status_t](#), [60](#)
 - [io_bw_encoding](#), [60](#)
- Energy Monitor (RAPL MSR), [16](#)
 - [esmi_all_energies_get](#), [17](#)
 - [esmi_core_energy_get](#), [16](#)
 - [esmi_socket_energy_get](#), [16](#)
- [esmi_all_energies_get](#)
 - Energy Monitor (RAPL MSR), [17](#)
- [esmi_apb_disable](#)
 - APB and LCLK level control, [40](#)
- [esmi_apb_enable](#)
 - APB and LCLK level control, [39](#)
- [esmi_cclk_limit_get](#)
 - HSMP System Statistics, [19](#)
- [esmi_core_boostlimit_get](#)
 - Performance (Boost limit) Monitor, [29](#)
- [esmi_core_boostlimit_set](#)
 - Performance (Boost limit) Control, [31](#)
- [esmi_core_energy_get](#)
 - Energy Monitor (RAPL MSR), [16](#)
- [esmi_cpu_family_get](#)
 - Auxiliary functions, [49](#)
- [esmi_cpu_model_get](#)
 - Auxiliary functions, [49](#)
- [esmi_current_freq_limit_core_get](#)
 - HSMP System Statistics, [21](#)
- [esmi_current_io_bandwidth_get](#)
 - Bandwidth Monitor, [43](#)
- [esmi_current_xgmi_bw_get](#)
 - Bandwidth Monitor, [43](#)
- [esmi_ddr_bw_get](#)
 - ddr_bandwidth Monitor, [33](#)
- [esmi_df_pstate_range_set](#)
 - APB and LCLK level control, [42](#)
- [esmi_dimm_power_consumption_get](#)
 - Dimm statistics, [35](#)
- [esmi_dimm_temp_range_and_refresh_rate_get](#)
 - Dimm statistics, [35](#)
- [esmi_dimm_thermal_sensor_get](#)
 - Dimm statistics, [36](#)
- [esmi_dram_address_metrics_table_get](#)
 - Metrics Table, [47](#)
- [esmi_fclk_mclk_get](#)
 - HSMP System Statistics, [19](#)
- [esmi_first_online_core_on_socket](#)
 - Auxiliary functions, [51](#)
- [esmi_get_err_msg](#)
 - Auxiliary functions, [51](#)
- [esmi_gmi3_link_width_range_set](#)
 - GMI3 width control, [38](#)
- [esmi_hsmp_proto_ver_get](#)
 - HSMP System Statistics, [20](#)
- [esmi_init](#)
 - Initialization and Shutdown, [15](#)
- [esmi_metrics_table_get](#)
 - Metrics Table, [45](#)
- [esmi_metrics_table_version_get](#)
 - Metrics Table, [45](#)
- [esmi_number_of_cpus_get](#)
 - Auxiliary functions, [50](#)
- [esmi_number_of_sockets_get](#)
 - Auxiliary functions, [51](#)
- [esmi_pcie_link_rate_set](#)
 - APB and LCLK level control, [41](#)
- [esmi_prochot_status_get](#)
 - HSMP System Statistics, [19](#)

- esmi_pwr_efficiency_mode_set
 - Power Control, [27](#)
- esmi_pwr_svi_telemetry_all_rails_get
 - Power Monitor, [24](#)
- esmi_smu_fw_version_get
 - HSMP System Statistics, [18](#)
- esmi_socket_boostlimit_set
 - Performance (Boost limit) Control, [31](#)
- esmi_socket_c0_residency_get
 - Performance (Boost limit) Monitor, [29](#)
- esmi_socket_current_active_freq_limit_get
 - HSMP System Statistics, [20](#)
- esmi_socket_energy_get
 - Energy Monitor (RAPL MSR), [16](#)
- esmi_socket_freq_range_get
 - HSMP System Statistics, [21](#)
- esmi_socket_lclk_dpm_level_get
 - APB and LCLK level control, [41](#)
- esmi_socket_lclk_dpm_level_set
 - APB and LCLK level control, [40](#)
- esmi_socket_power_cap_get
 - Power Monitor, [23](#)
- esmi_socket_power_cap_max_get
 - Power Monitor, [24](#)
- esmi_socket_power_cap_set
 - Power Control, [27](#)
- esmi_socket_power_get
 - Power Monitor, [23](#)
- esmi_socket_temperature_get
 - Temperature Query, [34](#)
- esmi_status_t
 - e_smi.h, [60](#)
- esmi_test_hsmp_mailbox
 - Test HSMP mailbox, [48](#)
- esmi_threads_per_core_get
 - Auxiliary functions, [50](#)
- esmi_xgmi_width_set
 - xGMI bandwidth control, [37](#)
- GMI3 width control, [38](#)
 - esmi_gmi3_link_width_range_set, [38](#)
- HSMP System Statistics, [18](#)
 - esmi_cclk_limit_get, [19](#)
 - esmi_current_freq_limit_core_get, [21](#)
 - esmi_fclk_mclk_get, [19](#)
 - esmi_hsmp_proto_ver_get, [20](#)
 - esmi_prochot_status_get, [19](#)
 - esmi_smu_fw_version_get, [18](#)
 - esmi_socket_current_active_freq_limit_get, [20](#)
 - esmi_socket_freq_range_get, [21](#)
- Initialization and Shutdown, [15](#)
 - esmi_init, [15](#)
- io_bw_encoding
 - e_smi.h, [60](#)
- link_id_bw_type, [55](#)
- Metrics Table, [45](#)
 - esmi_dram_address_metrics_table_get, [47](#)
 - esmi_metrics_table_get, [45](#)
 - esmi_metrics_table_version_get, [45](#)
- Performance (Boost limit) Control, [31](#)
 - esmi_core_boostlimit_set, [31](#)
 - esmi_socket_boostlimit_set, [31](#)
- Performance (Boost limit) Monitor, [29](#)
 - esmi_core_boostlimit_get, [29](#)
 - esmi_socket_c0_residency_get, [29](#)
- Power Control, [27](#)
 - esmi_pwr_efficiency_mode_set, [27](#)
 - esmi_socket_power_cap_set, [27](#)
- Power Monitor, [23](#)
 - esmi_pwr_svi_telemetry_all_rails_get, [24](#)
 - esmi_socket_power_cap_get, [23](#)
 - esmi_socket_power_cap_max_get, [24](#)
 - esmi_socket_power_get, [23](#)
- smu_fw_version, [55](#)
- temp_range_refresh_rate, [56](#)
- Temperature Query, [34](#)
 - esmi_socket_temperature_get, [34](#)
- Test HSMP mailbox, [48](#)
 - esmi_test_hsmp_mailbox, [48](#)
- xGMI bandwidth control, [37](#)
 - esmi_xgmi_width_set, [37](#)